

VIXL

A runtime assembler C++ API

Agenda

- What is VIXL?
- What is it used for?
- What's new in VIXL?
- Conclusion & Q&A

What is VIXL

- A Runtime code-generation library
 - for ARM AArch64 now, AArch32 (Arm and Thumb-2) coming CY2016
 - A program driven assembler, macro assembler, disassembler, simulator and debugger
- Developed as the code generator for the Google V8 JavaScript AArch64 port
- Available through GitHub
 - <https://github.com/armvixl/vixl>
- Supports most of the instruction set
 - At ELI (usertime)
 - With full NEON support
- Used by a number of projects
 - Android's ART, QEMU, HipHop, Mozilla, and more!

VIXL's Assembler

- Example of a loop:

```
    mov x0, 10
loop: subs x0, x0, 1
      bne loop
```

- VIXL

```
Label loop;
mov(x0, 10);
Bind(&loop);
subs(x0, x0, 1);
b(&loop, ne);
```

NEON support in the assembler

- Example of generating NEON code:
`Ld1(v0.V8B, v1.V8B(), MemOperand(x0));`
`Sadd1(v0.V8H(), v0.V8B(), v1.V8B());`
`Addv(h0, v0.V8H());`
- Will produce this code:
`ld1 {v0.8b, v1.8b}, [x0]`
`sadd1 v0.8h, v0.8b, v1.8b`
`addv h0, v0.8h`
- Register methods (here, for example `V8H()`) determine
 - the number of lanes (8)
 - and size of the register (D or Q) ($8 * 16\text{bit} = 128\text{bits} = \text{Q}$)

VIXL's MacroAssembler

- While VIXL's Assembler will generate one instruction, the MacroAssembler gives you a lot more flexibility
 - AArch64 uses modified immediates
 - `mov r0, 0x10001` is permitted
 - `mov r0, 0x10002` is not!
 - `MacroAssembler::Mov(r0, 0x10002)` will do the job for you, using a pair of `movz/movk` or whatever combination of instructions necessary
 - The MacroAssembler ultimately calls the Assembler
- Simplifies your code without a deep knowledge of ARM AArch64
- Sometimes it might optimize your code
 - But it's still an assembler, not a compiler!

VIXL's MacroAssembler (2)

- Will perform simple tricks to shorten your generated code

- For example:

```
Add(w0, w0, 0x10002);
```

- Will generate:

```
mov w16, #0x8001
```

```
add x0, x0, w16, lsl #1
```

- Instead of the usual movz/movk for 0x10002

What to do with the generated code

- In a JIT environment, simply run it
 - Natively,
 - Or Simulated
- Save it for future reuse
- Or generate an ELF file using elfio and link it to your application
 - Create a symbol for the generated functions using elfio

```
Elf_word add_symbol(string_section_accessor& pStrWriter,
                    const char* str, Elf64_Addr value,
                    Elf_Xword size, unsigned char info,
                    unsigned char other, Elf_Half shndx );
symbols.add_symbol( strings, fn_name, fn_addr, fn_size, STB_GLOBAL,
                    STB_FUNC, 0, text_sec->get_index);
```
 - ELFio is available here: <http://elfio.sourceforge.net>

A quick example

- The headers

```
#include <stdio.h>
#include "vixl/a64/assembler-a64.h"
#include "vixl/a64/disasm-a64.h"

using namespace vixl;
```

A quick example (cont ...)

```
int main(int argc, char ** argv) {
    const size_t bufSize = 1024;
    uint8_t* buffer = new uint8_t[bufSize];
    Assembler asm_(buffer, bufSize);

    asm_.orr(w0, wzr, 0x10001);
    asm_.mov(x1, x0);
    asm_.FinalizeCode();

    BufferDisassembler disasm_(buffer,
        asm_.CursorOffset());
}
```

A quick example (cont ...)

```
class BufferDisassembler {
    Decoder decoder_;
    PrintDisassembler disasm_;
public:
    BufferDisassembler(
        uint8_t* buffer, size_t size)
        : disasm_(stdout) {
        decoder_.AppendVisitor(&disasm_);
        uint32_t* instr =
            reinterpret_cast<uint32_t*>(buffer);
        for (; size > 0; size -= 4, instr++) {
            decoder_.Decode(
                reinterpret_cast<Instruction*>(instr));
        }
    }
};
```

//

A quick example (cont ...)

- Will output

```
0x00000000013ad010  320083e0  mov w0, #0x10001
0x00000000013ad014  aa0003e1  mov x1, x0
```

Another one using the macro assembler

```
asm_.Mov(w0, 0x10002);  
asm_.Mov(x1, x0);
```

- Will output:

0x0000000000f3f010	52800040	mov w0, #0x2
0x0000000000f3f014	72a00020	movk w0, #0x1, lsl #16
0x0000000000f3f018	aa0003e1	mov x1, x0

- A lot more examples (and surely more complex) are available in the example directory

VIXL's Simulator

```
0x000000002ff90b4 4c960a40          st4 {v0.4s, v1.4s, v2.4s, v3.4s}, [x18], x22
#   v0: 0x0f0e0d0c0b0a09080706050403020100 (7.00365e-30, 2.65846e-32, 1.00825e-34, 3.82047e-37) <- 0x00007ffd447edf15
#   v1: 0x1f1eld1c1b1a19181716151413121110 (3.34819e-20, 1.27467e-22, 4.84942e-25, 1.84362e-27) <- 0x00007ffd447edf15
#   v2: 0x2f2e2d2c2b2a29282726252423222120 (1.58413e-10, 6.04532e-13, 2.30573e-15, 8.78905e-18) <- 0x00007ffd447edf15
#   v3: 0x3f3e3d3c3b3a39383736353433323130 (0.743122, 0.00284155, 1.08604e-05, 4.14886e-08) <- 0x00007ffd447edf15
#   x18: 0x00007ffd447edf1a
0x000000002ff90b8 4c000e40          st4 {v0.2d, v1.2d, v2.2d, v3.2d}, [x18]
#   v0: 0x0f0e0d0c0b0a09080706050403020100 (3.69192e-236, 7.94993e-275) <- 0x00007ffd447edf1a
#   v1: 0x1f1eld1c1b1a19181716151413121110 (8.56775e-159, 1.84632e-197) <- 0x00007ffd447edf1a
#   v2: 0x2f2e2d2c2b2a29282726252423222120 (1.98829e-81, 4.28794e-120) <- 0x00007ffd447edf1a
#   v3: 0x3f3e3d3c3b3a39383736353433323130 (0.000461414, 9.95833e-43) <- 0x00007ffd447edf1a
0x000000002ff90bc d29bda13          mov x19, #0xded0
#   x19: 0x00000000000000ded0
0x000000002ff90c0 f2a88fd3          movk x19, #0x447e, lsl #16
#   x19: 0x000000000447eded0
0x000000002ff90c4 f2cffffb3          movk x19, #0x7ffd, lsl #32
#   x19: 0x00007ffd447eded0
0x000000002ff90c8 3cc10660          ldr q0, [x19], #16
#   v0: 0x12030231302120111001000130201000 <- 0x00007ffd447eded0
#   x19: 0x00007ffd447edee0
0x000000002ff90cc 3cc10661          ldr q1, [x19], #16
#   v1: 0x16070635342524151405043332232213 <- 0x00007ffd447edee0
#   x19: 0x00007ffd447edef0
0x000000002ff90d0 3cc10662          ldr q2, [x19], #16
#   v2: 0x1a0b0a39382928191809083736272617 <- 0x00007ffd447edef0
#   x19: 0x00007ffd447edf00
```

- Able to simulate whatever instruction VIXL can assemble
- Execution and register trace output available
- Extensive back-to-back testing against hardware
- Printf!

VIXL Simulator (cont)

- Because everything is simulated
 - Every instruction executed is traced
 - Every register update is traced
 - Every memory load/store, with address information, is checked
- Can be seen as slow
 - No, it's not! if run on many cores, at let's say 3 or 4Ghz, it will be fast!

VIXL's Debugger

- gdb with fewer commands
 - p for print, x for examine, si to execute the current instruction

```
MacroAssembler masm(assm_buf, BUF_SIZE);
Decoder decoder;
Debugger debugger(&decoder);

...
// Generate the code for the example function.
Label start;
masm.Bind(&start);
masm.Brk();
masm.Nop();
. . .
debugger.RunFrom(
    masm.GetLabelAddress<Instruction*>(&start));
```

```
Hit breakpoint at pc=0x7ffe93ce2a40.
Next: 0x00007ffe93ce2a40 d4200000 brk #0x0
vixl> si
Next: 0x00007ffe93ce2a48 14000002 b #+0x8 (addr
0x7ffe93ce2a50)
vixl>
Next: 0x00007ffe93ce2a50 d2800f61 mov x1, #0x7b
vixl>
Next: 0x00007ffe93ce2a54 d2803902 mov x2, #0x1c8
vixl>
Next: 0x00007ffe93ce2a58 8b020020 add x0, x1, x2
vixl>
Next: 0x00007ffe93ce2a5c d65f03c0 ret
vixl> p x1
x1 = 0000000000000007b
vixl> p x2
x2 = 000000000000001c8
vixl> p x0
x0 = 0000000000000243
vixl> c
```


Decoders

- The disassembler and simulator take action based on a common instruction decoder
- They use a visitor pattern
 - For each instruction, call the visitor methods of all registered classes
 - VisitAddSubImmediate()
 - VisitDataProcessing3Source()
- Classes associated with instruction decoding implement this visitor interface
 - Simulator
 - Disassembler
 - Instrumentation
- Add your own visitors to get reports on what instructions have been encountered

Hacking with VIXL

- Most projects only use part of VIXL
 - Android's ART uses the assembler and the disassembler
 - QEMU uses the disassembler
- Google V8 and Mozilla use everything
 - The simulator is used to run tests on many fast cores
- The components are useful in isolation
 - We have already produced small command line tools for convenience, like a64disasm

VIXL as a programmable assembler

- Use VIXL to generate static code from C++, produce an object, and link it into the final executable
 - Plug VIXL into elfio to make a programmable assembler
 - Like an inline assembler, but more flexible
 - Like intrinsics, but for all instructions
- Example of an inner loop of an image processing operation:

```
Bind(&loop);  
Ldrb(value, MemOperand(data));  
Mul(value, value, coeff);  
Lsr(value, value, 8);  
Strb(value, MemOperand(data, 1, PostIndex));  
Sub(length, length, 1);  
Cbnz(length, &loop);
```

VIXL as a programmable assembler (2)

- Generate the code for different values of `coeff`
- Same code sequence can produce special cases for each value
- Create a symbol for each of the generated functions using elfio

```
symbols.add_symbol(strings, fn_name, fn_addr, fn_size, STB_GLOBAL,  
                    STB_FUNC, 0, text_sec->get_index);
```
- Build the programmable assembler
- Run it to produce objects that can be linked into the rest of your project
- Like GAS, but more flexible
- Or wait from inside your app, generate the function upon request, and call it
 - You have created a dynamic C++ template instantiation!

```
template <int8_t coeff> mullcoeff(int8_t *data);
```

Dynamic coding example

```
Function mul(MacroAssembler& asm_, uint8_t coeff) {  
    ptrdiff_t fn_start= asm_.CursorOffset();  
    Label mul;  
    asm_.Bind(&mul);  
    Label loop;  
    asm_.Bind(&loop);  
    asm_.Ldrb(w11, MemOperand(x0));  
    asm_.Mul(w11, w11, w1);  
    asm_.Lsr(w11, w11, 8);  
    asm_.Strb(w11, MemOperand(x0, 1, PostIndex));  
    asm_.Ret();  
    return Function(mul.location(), asm_.CursorOffset() - fn_start);  
}
```

Dynamic coding example (cont...)

- The obvious cases

```
if (coeff != 1) {
    Label loop;
    asm_.Bind(&loop);
    if (coeff == 0) {
        asm_.Strb(wzr, MemOperand(x0, 1, PostIndex));
        asm_.Sub(x1, x1, 1);
        asm_.Cbnz(x1, &loop);
    } else {
        ...
    }
}
asm_.Ret();
```

Dynamic coding example (cont...)

mul_0:

0x1fd2010 3800141f

0x1fd2014 d1000421

0x1fd2018 b5ffffc1

0x1fd201c d65f03c0

mul_1:

0x1fd2020 d65f03c0

strb wzr, [x0], #1

sub x1, x1, #0x1 (1)

cbnz x1, #-0x8 (addr 0x1fd2010)

ret

ret

Dynamic coding example (cont...)

- What if coeff is a power of 2?

```
Label loop;  
asm_.Bind(&loop);  
asm_.Ldrb(w11, MemOperand(x0));  
if (IsPowerOf2(coeff)) {  
    unsigned pow2= CountTrailingZeros(coeff);  
    asm_.Ubfx(w11, w11, 8 - pow2, pow2);  
} else {  
    ...  
}
```


Dynamic coding example (cont...)

mul_2:

0x1fd2024	3940000b	ldrb w11, [x0]
0x1fd2028	53077d6b	ubfx w11, w11, #7, #1
0x1fd202c	3800140b	strb w11, [x0], #1
0x1fd2030	d1000421	sub x1, x1, #0x1 (1)
0x1fd2034	b5ffff81	cbnz x1, #-0x10 (addr 0x1fd2024)
0x1fd2038	d65f03c0	ret

Dynamic coding example (cont...)

```
int main(int argc, char ** argv) {  
    Function jumptable[255];  
    const size_t bufSize = 1024;  
    uint8_t* buffer = new uint8_t[bufSize];  
    MacroAssembler asm_(buffer, bufSize);  
  
    jumptable[0] = mul(asm_, 0);  
    jumptable[1] = mul(asm_, 1);  
    jumptable[2] = mul(asm_, 2);  
    jumptable[11] = mul(asm_, 11);  
  
    asm_.FinalizeCode();  
}
```

- jumpTable can also be an std::set and upper_bound can be used
 - Because jumpTable[11] is the same as jumpTable[12]

Random instruction set testing

- To test your CPU by brute force
 - Execute a constrained random instruction stream, and see what breaks
- Use VIXL to generate this on the target under test rather than ahead of time
- Generate random 32bit values, and use VIXL to decide if the instruction is valid before executing natively
- Compare observable state after an instruction sequence between simulator and native, and detect deviation
 - The simulator can be compiled for AArch64, so simulated and native instructions can be run in the same environment
- Use simulator feedback to inform decisions about what to generate
 - Check address register is legal for generating memory access ops
 - Verify loops can exit in a reasonable amount of time

Where to use VIXL

- JITs: JavaScript, Java, Python, other scripting languages
- Dynamic code generation of optimized routine
- Testing
 - Random Instruction Stream (RIS) testing
 - Toolchain testing
- ISA experimentation: try out features of ARM's AArch64
 - Simple, fast, tested API
 - Integrated suite, ready to use on a new JIT project
- Teaching tool using a parsing expression grammar tool like PEG.js
 - Turn this JSON representation
 - {“mnemonic”:”mov”, “destination”: {“type”:”x”, “index”:l}, ...
 - Into a Mov(xl, ...) can be done in a matter of hours

What's new in VIXL

- VIXL1.8 introduced a few more things:
 - Full NEON support
 - Support for long branches
 - Improved handling of literal pools.
 - Support some `ic` and `dc` cache op instructions.
 - Support CRC32 instructions.
 - Support half-precision floating point instructions.
 - MacroAssembler support for `bfm`, `ubfm` and `sbfm`.
 - Other small bug fixes and improvements.
- We are now at 1.10 (July 2015)
 - improved literal handling,
 - a better build and test infrastructure,
 - And bug fixes

Conclusion

- Vixl is a toolkit
 - VIXL does the boring bits for you
- VIXL is available on github
 - <https://github.com/armvixl/vixl>
 - The license is a permissive free software license equivalent to the BSD 3-Clause license
- VIXL is maintained and developed
 - We are now (since July 2015) at 1.10
 - Supported by ARM
- ... and extended
 - VIXL (= LXIV, ie 64 in roman numerals) for Aarch32 is coming CY2016
- A good reading start:
 - <https://github.com/armvixl/vixl/blob/master/README.md>